

# phw3

May 11, 2025

```
[2]: # Necessary libraries
import h5py
import numpy as np
import matplotlib.pyplot as plt
import math
from collections import Counter
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.utils import shuffle
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

from tensorflow.keras.utils import to_categorical

import librosa, librosa.display, numpy as np, matplotlib.pyplot as plt, pandas_
↪as pd
from tensorflow.keras.models import load_model
import pathlib
```

```
[3]: # loading the file
data = h5py.File('birds/bird_spectrograms.hdf5', 'r')

bird_names = list(data.keys())
print(bird_names)

#labeling them for readability
names = {
    'amecro': 'American Crow',
    'amerob': 'American Robin',
    'bewwre': 'Bewicks Wren',
    'bkcchi': 'Black capped Chickadee',
    'daejun': 'Dark eyed Junco',
    'houfin': 'House Finch',
    'houspa': 'House Sparrow',
    'norfli': 'Northern Flicker',
    'rewbla': 'Red winged Blackbird',
    'sonspa': 'Song Sparrow',
```

```

        'spotow': 'Spotted Towhee',
        'whcspa': 'White crowned Sparrow'
    }
    print(names)

    print("All bird species:")
    all_birds = [names[code] for code in bird_names]
    for bird in all_birds:
        print(bird)

    for code in bird_names:
        shape = data[code].shape
        print(f"{code:7s} {names.get(code, 'Unknown')}:25s} {shape}")

    n_cols, n_rows = 3, math.ceil(len(bird_names) / 3)
    fig, axes = plt.subplots(n_rows, n_cols, figsize=(12, 4 * n_rows))
    axes = axes.flatten()

    for i, code in enumerate(bird_names):
        spectro = data[code][:, :, 0]
        axes[i].imshow(spectro, aspect="auto", origin="lower", cmap="gray")
        axes[i].set_title(f"{code} | {names.get(code, '')}", fontsize=9)
        axes[i].set_xlabel("Time bins")
        axes[i].set_ylabel("Frequency bins")

    for j in range(i + 1, len(axes)):
        axes[j].axis("off")

    plt.tight_layout()
    plt.show()

```

```

['amecro', 'amerob', 'bewwre', 'bkcchi', 'daejun', 'houfin', 'houspa', 'norfli',
'rewbla', 'sonspa', 'spotow', 'whcspa']
{'amecro': 'American Crow', 'amerob': 'American Robin', 'bewwre': 'Bewicks
Wren', 'bkcchi': 'Black capped Chickadee', 'daejun': 'Dark eyed Junco',
'houfin': 'House Finch', 'houspa': 'House Sparrow', 'norfli': 'Northern
Flicker', 'rewbla': 'Red winged Blackbird', 'sonspa': 'Song Sparrow', 'spotow':
'Spotted Towhee', 'whcspa': 'White crowned Sparrow'}

```

All bird species:

American Crow

American Robin

Bewicks Wren

Black capped Chickadee

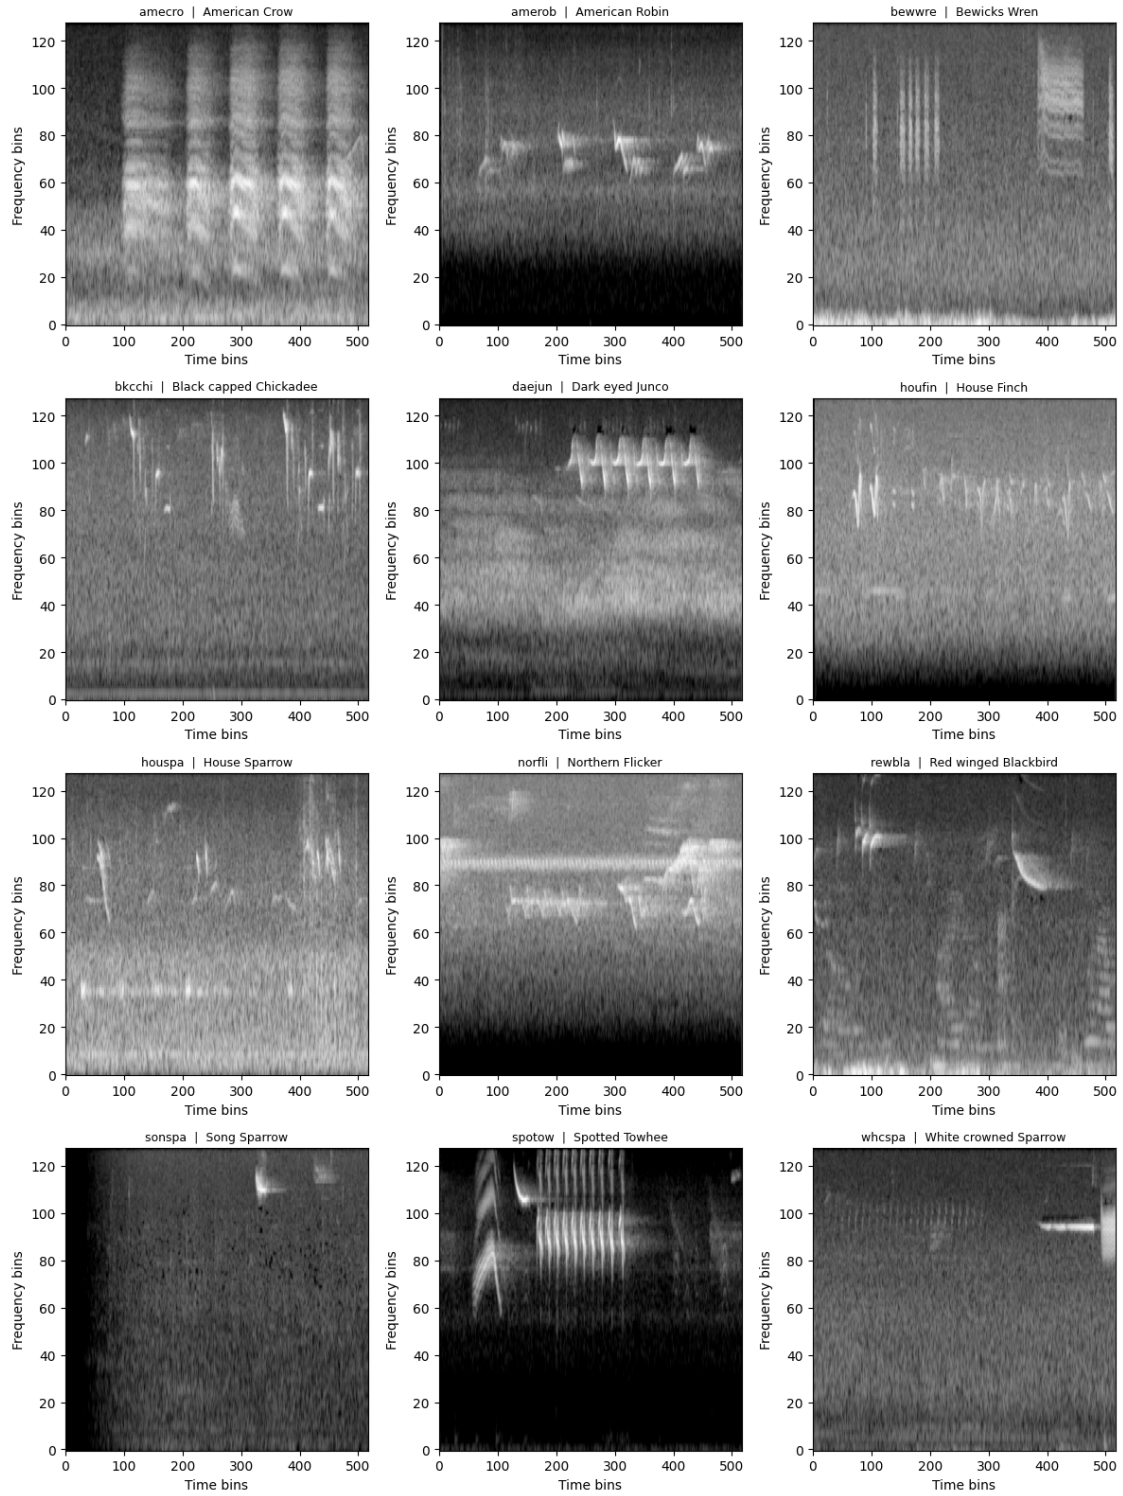
Dark eyed Junco

House Finch

House Sparrow

Northern Flicker

Red winged Blackbird  
 Song Sparrow  
 Spotted Towhee  
 White crowned Sparrow  
 amecro American Crow (128, 517, 66)  
 amerob American Robin (128, 517, 172)  
 bewwre Bewicks Wren (128, 517, 144)  
 bkcchi Black capped Chickadee (128, 517, 45)  
 daejun Dark eyed Junco (128, 517, 125)  
 houfin House Finch (128, 517, 84)  
 houspa House Sparrow (128, 517, 630)  
 norfli Northern Flicker (128, 517, 37)  
 rewbla Red winged Blackbird (128, 517, 187)  
 sonspa Song Sparrow (128, 517, 263)  
 spotow Spotted Towhee (128, 517, 137)  
 whcspa White crowned Sparrow (128, 517, 91)



```
[4]: with data as f:
      song_sparrow = f['sonspa'][...].astype(np.float32) / 255.0
```

```

white_sparrow = f['whcspa'][...].astype(np.float32) / 255.0

# making into 2-d
song_sparrow = np.transpose(song_sparrow, (2, 0, 1))
white_sparrow = np.transpose(white_sparrow, (2, 0, 1))

# making both the classes the same length
minimum_length = min(song_sparrow.shape[0], white_sparrow.shape[0])
song_sparrow = song_sparrow[:minimum_length]
white_sparrow = white_sparrow[:minimum_length]

# Labels, 0 for Song Sparrow, 1 for white crowned sparrow
song_one = np.zeros(minimum_length, dtype=np.uint8)
white_one = np.ones(minimum_length, dtype=np.uint8)

x_class = np.concatenate([song_sparrow, white_sparrow], axis=0)
y_class = np.concatenate([song_one, white_one], axis=0)
x_class = x_class[..., np.newaxis]
x_class, y_class = shuffle(x_class, y_class, random_state=42)

# Splitting 30% train, 40% test, 30% validation set
x_one, x_test, y_one, y_test = train_test_split(x_class, y_class,
↳test_size=0.4, stratify=y_class, random_state=42)
x_train, x_valid, y_train, y_valid = train_test_split(x_one, y_one,
↳test_size=0.5, stratify=y_one, random_state=42)

```

```

[ ]: # binary CNN model
model_one = tf.keras.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu', input_shape=x_train.
↳shape[1:]),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dropout(0.3),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

model_one.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-4),
    loss='binary_crossentropy',
    metrics=['accuracy', tf.keras.metrics.AUC(name='auc')]
)

```

```

model_one.summary()

callbacks_one = [
    tf.keras.callbacks.EarlyStopping(patience=10, restore_best_weights=True),
    tf.keras.callbacks.ModelCheckpoint("lecture_cnn_sparrows.keras",
    ↪save_best_only=True)
]

history = model_one.fit(
    x_train, y_train,
    validation_data=(x_valid, y_valid),
    epochs=100,
    batch_size=32,
    callbacks=callbacks_one,
    verbose=1
)

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.legend()
plt.title('Accuracy')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.title('Loss')

plt.tight_layout()
plt.show()

loss, acc, auc = model_one.evaluate(x_test, y_test, verbose=0)
print(f"\n Final Test Accuracy: {acc:.4f} | AUC: {auc:.4f}")

probability_one = model_one.predict(x_test)
prediction_one = (probability_one > 0.5).astype(int).flatten()

conf_matrix = confusion_matrix(y_test, prediction_one)
matrix_one = ConfusionMatrixDisplay(confusion_matrix=conf_matrix,
    ↪display_labels=['Song Sparrow', 'White crowned Sparrow'])
matrix_one.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 515, 16)	160
max_pooling2d (MaxPooling2D)	(None, 63, 257, 16)	0
conv2d_1 (Conv2D)	(None, 61, 255, 31)	4495
max_pooling2d_1 (MaxPooling2D)	(None, 30, 127, 31)	0
flatten (Flatten)	(None, 118110)	0
dense (Dense)	(None, 32)	3779552
dropout (Dropout)	(None, 32)	0
dense_1 (Dense)	(None, 1)	33

```

=====
Total params: 3,784,240
Trainable params: 3,784,240
Non-trainable params: 0
=====

```

```

-----
Epoch 1/100
2/2 [=====] - 1s 645ms/step - loss: 0.6888 - accuracy:
0.5000 - auc: 0.6187 - val_loss: 0.6864 - val_accuracy: 0.4909 - val_auc: 0.7698
Epoch 2/100
2/2 [=====] - 1s 438ms/step - loss: 0.7004 - accuracy:
0.5000 - auc: 0.5123 - val_loss: 0.6800 - val_accuracy: 0.4909 - val_auc: 0.7884
Epoch 3/100
2/2 [=====] - 1s 454ms/step - loss: 0.6759 - accuracy:
0.5926 - auc: 0.7154 - val_loss: 0.6774 - val_accuracy: 0.6182 - val_auc: 0.7930
Epoch 4/100
2/2 [=====] - 1s 421ms/step - loss: 0.6854 - accuracy:
0.6111 - auc: 0.6043 - val_loss: 0.6749 - val_accuracy: 0.6182 - val_auc: 0.7884
Epoch 5/100
2/2 [=====] - 1s 390ms/step - loss: 0.6797 - accuracy:
0.6296 - auc: 0.6529 - val_loss: 0.6686 - val_accuracy: 0.5455 - val_auc: 0.7844
Epoch 6/100
2/2 [=====] - 1s 397ms/step - loss: 0.6613 - accuracy:
0.6481 - auc: 0.7641 - val_loss: 0.6662 - val_accuracy: 0.5091 - val_auc: 0.7844
Epoch 7/100
2/2 [=====] - 1s 400ms/step - loss: 0.6587 - accuracy:
0.6111 - auc: 0.7867 - val_loss: 0.6628 - val_accuracy: 0.5455 - val_auc: 0.7890
Epoch 8/100

```

2/2 [=====] - 1s 410ms/step - loss: 0.6529 - accuracy:  
0.6296 - auc: 0.7538 - val\_loss: 0.6587 - val\_accuracy: 0.6182 - val\_auc: 0.7864  
Epoch 9/100

2/2 [=====] - 1s 394ms/step - loss: 0.6497 - accuracy:  
0.7037 - auc: 0.7682 - val\_loss: 0.6555 - val\_accuracy: 0.6909 - val\_auc: 0.7831  
Epoch 10/100

2/2 [=====] - 1s 414ms/step - loss: 0.6703 - accuracy:  
0.6111 - auc: 0.6584 - val\_loss: 0.6532 - val\_accuracy: 0.7273 - val\_auc: 0.7890  
Epoch 11/100

2/2 [=====] - 1s 398ms/step - loss: 0.6486 - accuracy:  
0.7222 - auc: 0.7483 - val\_loss: 0.6504 - val\_accuracy: 0.6909 - val\_auc: 0.7831  
Epoch 12/100

2/2 [=====] - 1s 396ms/step - loss: 0.6355 - accuracy:  
0.7222 - auc: 0.7846 - val\_loss: 0.6485 - val\_accuracy: 0.6909 - val\_auc: 0.7831  
Epoch 13/100

2/2 [=====] - 1s 411ms/step - loss: 0.6470 - accuracy:  
0.6852 - auc: 0.7250 - val\_loss: 0.6466 - val\_accuracy: 0.7091 - val\_auc: 0.7851  
Epoch 14/100

2/2 [=====] - 1s 442ms/step - loss: 0.6441 - accuracy:  
0.5926 - auc: 0.7222 - val\_loss: 0.6444 - val\_accuracy: 0.7091 - val\_auc: 0.7791  
Epoch 15/100

2/2 [=====] - 1s 560ms/step - loss: 0.6499 - accuracy:  
0.6481 - auc: 0.7023 - val\_loss: 0.6441 - val\_accuracy: 0.6727 - val\_auc: 0.7665  
Epoch 16/100

2/2 [=====] - 1s 613ms/step - loss: 0.6519 - accuracy:  
0.6667 - auc: 0.7119 - val\_loss: 0.6405 - val\_accuracy: 0.6727 - val\_auc: 0.7685  
Epoch 17/100

2/2 [=====] - 1s 547ms/step - loss: 0.6392 - accuracy:  
0.7037 - auc: 0.7332 - val\_loss: 0.6362 - val\_accuracy: 0.6727 - val\_auc: 0.7738  
Epoch 18/100

2/2 [=====] - 1s 600ms/step - loss: 0.6227 - accuracy:  
0.7037 - auc: 0.7545 - val\_loss: 0.6318 - val\_accuracy: 0.7273 - val\_auc: 0.7778  
Epoch 19/100

2/2 [=====] - 1s 568ms/step - loss: 0.6107 - accuracy:  
0.7222 - auc: 0.7840 - val\_loss: 0.6283 - val\_accuracy: 0.7091 - val\_auc: 0.7817  
Epoch 20/100

2/2 [=====] - 1s 600ms/step - loss: 0.5914 - accuracy:  
0.7407 - auc: 0.8114 - val\_loss: 0.6267 - val\_accuracy: 0.7273 - val\_auc: 0.7864  
Epoch 21/100

2/2 [=====] - 1s 567ms/step - loss: 0.6107 - accuracy:  
0.6481 - auc: 0.7798 - val\_loss: 0.6231 - val\_accuracy: 0.7273 - val\_auc: 0.7857  
Epoch 22/100

2/2 [=====] - 1s 571ms/step - loss: 0.6011 - accuracy:  
0.7222 - auc: 0.7819 - val\_loss: 0.6200 - val\_accuracy: 0.6909 - val\_auc: 0.7837  
Epoch 23/100

2/2 [=====] - 1s 580ms/step - loss: 0.6295 - accuracy:  
0.6296 - auc: 0.7318 - val\_loss: 0.6190 - val\_accuracy: 0.6909 - val\_auc: 0.7837  
Epoch 24/100



2/2 [=====] - 1s 567ms/step - loss: 0.6742 - accuracy: 0.6111 - auc: 0.6255 - val\_loss: 0.6190 - val\_accuracy: 0.6727 - val\_auc: 0.7791  
Epoch 25/100

2/2 [=====] - 1s 567ms/step - loss: 0.6538 - accuracy: 0.6296 - auc: 0.6818 - val\_loss: 0.6181 - val\_accuracy: 0.6727 - val\_auc: 0.7784  
Epoch 26/100

2/2 [=====] - 1s 550ms/step - loss: 0.6661 - accuracy: 0.6296 - auc: 0.6536 - val\_loss: 0.6133 - val\_accuracy: 0.6909 - val\_auc: 0.7870  
Epoch 27/100

2/2 [=====] - 1s 562ms/step - loss: 0.6168 - accuracy: 0.7407 - auc: 0.7373 - val\_loss: 0.6083 - val\_accuracy: 0.6909 - val\_auc: 0.7917  
Epoch 28/100

2/2 [=====] - 1s 561ms/step - loss: 0.5789 - accuracy: 0.7407 - auc: 0.8045 - val\_loss: 0.6064 - val\_accuracy: 0.6909 - val\_auc: 0.7903  
Epoch 29/100

2/2 [=====] - 1s 500ms/step - loss: 0.6017 - accuracy: 0.7222 - auc: 0.7730 - val\_loss: 0.6068 - val\_accuracy: 0.7091 - val\_auc: 0.7917  
Epoch 30/100

2/2 [=====] - 1s 566ms/step - loss: 0.6041 - accuracy: 0.7222 - auc: 0.7778 - val\_loss: 0.6060 - val\_accuracy: 0.7091 - val\_auc: 0.7903  
Epoch 31/100

2/2 [=====] - 1s 570ms/step - loss: 0.5994 - accuracy: 0.7407 - auc: 0.7922 - val\_loss: 0.6027 - val\_accuracy: 0.6909 - val\_auc: 0.7950  
Epoch 32/100

2/2 [=====] - 1s 569ms/step - loss: 0.5988 - accuracy: 0.7037 - auc: 0.7599 - val\_loss: 0.5996 - val\_accuracy: 0.7091 - val\_auc: 0.7956  
Epoch 33/100

2/2 [=====] - 1s 550ms/step - loss: 0.5926 - accuracy: 0.7037 - auc: 0.7997 - val\_loss: 0.5955 - val\_accuracy: 0.7091 - val\_auc: 0.7970  
Epoch 34/100

2/2 [=====] - 1s 559ms/step - loss: 0.5669 - accuracy: 0.7778 - auc: 0.8182 - val\_loss: 0.5926 - val\_accuracy: 0.7273 - val\_auc: 0.7989  
Epoch 35/100

2/2 [=====] - 1s 583ms/step - loss: 0.5762 - accuracy: 0.7037 - auc: 0.8086 - val\_loss: 0.5905 - val\_accuracy: 0.7091 - val\_auc: 0.8009  
Epoch 36/100

2/2 [=====] - 1s 567ms/step - loss: 0.5846 - accuracy: 0.7593 - auc: 0.8059 - val\_loss: 0.5883 - val\_accuracy: 0.7273 - val\_auc: 0.8009  
Epoch 37/100

2/2 [=====] - 1s 558ms/step - loss: 0.5704 - accuracy: 0.6852 - auc: 0.8251 - val\_loss: 0.5852 - val\_accuracy: 0.7273 - val\_auc: 0.8049  
Epoch 38/100

2/2 [=====] - 1s 550ms/step - loss: 0.5699 - accuracy: 0.7593 - auc: 0.8141 - val\_loss: 0.5851 - val\_accuracy: 0.7091 - val\_auc: 0.8036  
Epoch 39/100

2/2 [=====] - 1s 484ms/step - loss: 0.5710 - accuracy: 0.7037 - auc: 0.8045 - val\_loss: 0.5876 - val\_accuracy: 0.6727 - val\_auc: 0.8029  
Epoch 40/100

2/2 [=====] - 1s 484ms/step - loss: 0.6029 - accuracy:  
0.7037 - auc: 0.7695 - val\_loss: 0.5858 - val\_accuracy: 0.7091 - val\_auc: 0.8022  
Epoch 41/100

2/2 [=====] - 1s 500ms/step - loss: 0.5892 - accuracy:  
0.7222 - auc: 0.7785 - val\_loss: 0.5856 - val\_accuracy: 0.7091 - val\_auc: 0.8036  
Epoch 42/100

2/2 [=====] - 1s 569ms/step - loss: 0.5294 - accuracy:  
0.7778 - auc: 0.8512 - val\_loss: 0.5823 - val\_accuracy: 0.7273 - val\_auc: 0.8056  
Epoch 43/100

2/2 [=====] - 1s 484ms/step - loss: 0.5513 - accuracy:  
0.7778 - auc: 0.8388 - val\_loss: 0.5785 - val\_accuracy: 0.7273 - val\_auc: 0.8069  
Epoch 44/100

2/2 [=====] - 1s 500ms/step - loss: 0.5244 - accuracy:  
0.7593 - auc: 0.8567 - val\_loss: 0.5739 - val\_accuracy: 0.7273 - val\_auc: 0.8069  
Epoch 45/100

2/2 [=====] - 1s 520ms/step - loss: 0.5671 - accuracy:  
0.7407 - auc: 0.7805 - val\_loss: 0.5707 - val\_accuracy: 0.7273 - val\_auc: 0.8082  
Epoch 46/100

2/2 [=====] - 1s 504ms/step - loss: 0.5658 - accuracy:  
0.6667 - auc: 0.7915 - val\_loss: 0.5706 - val\_accuracy: 0.7273 - val\_auc: 0.8075  
Epoch 47/100

2/2 [=====] - 1s 487ms/step - loss: 0.5485 - accuracy:  
0.6852 - auc: 0.8073 - val\_loss: 0.5695 - val\_accuracy: 0.7091 - val\_auc: 0.8102  
Epoch 48/100

2/2 [=====] - 1s 487ms/step - loss: 0.5396 - accuracy:  
0.7593 - auc: 0.8189 - val\_loss: 0.5656 - val\_accuracy: 0.7273 - val\_auc: 0.8102  
Epoch 49/100

2/2 [=====] - 1s 500ms/step - loss: 0.4921 - accuracy:  
0.7963 - auc: 0.8896 - val\_loss: 0.5611 - val\_accuracy: 0.7273 - val\_auc: 0.8095  
Epoch 50/100

2/2 [=====] - 1s 494ms/step - loss: 0.5410 - accuracy:  
0.7593 - auc: 0.8073 - val\_loss: 0.5598 - val\_accuracy: 0.7636 - val\_auc: 0.8102  
Epoch 51/100

2/2 [=====] - 1s 483ms/step - loss: 0.5063 - accuracy:  
0.7593 - auc: 0.8567 - val\_loss: 0.5569 - val\_accuracy: 0.7455 - val\_auc: 0.8082  
Epoch 52/100

2/2 [=====] - 1s 499ms/step - loss: 0.5449 - accuracy:  
0.7407 - auc: 0.7956 - val\_loss: 0.5563 - val\_accuracy: 0.7273 - val\_auc: 0.8082  
Epoch 53/100

2/2 [=====] - 1s 438ms/step - loss: 0.6184 - accuracy:  
0.6667 - auc: 0.7270 - val\_loss: 0.5647 - val\_accuracy: 0.7273 - val\_auc: 0.8102  
Epoch 54/100

2/2 [=====] - 1s 437ms/step - loss: 0.4712 - accuracy:  
0.8148 - auc: 0.8978 - val\_loss: 0.5664 - val\_accuracy: 0.7273 - val\_auc: 0.8102  
Epoch 55/100

2/2 [=====] - 1s 516ms/step - loss: 0.5126 - accuracy:  
0.7407 - auc: 0.8237 - val\_loss: 0.5546 - val\_accuracy: 0.7455 - val\_auc: 0.8089  
Epoch 56/100

2/2 [=====] - 1s 500ms/step - loss: 0.4885 - accuracy:  
0.7963 - auc: 0.8656 - val\_loss: 0.5511 - val\_accuracy: 0.7636 - val\_auc: 0.8115  
Epoch 57/100  
2/2 [=====] - 1s 498ms/step - loss: 0.5240 - accuracy:  
0.7593 - auc: 0.8381 - val\_loss: 0.5503 - val\_accuracy: 0.7636 - val\_auc: 0.8122  
Epoch 58/100  
2/2 [=====] - 1s 537ms/step - loss: 0.5209 - accuracy:  
0.7778 - auc: 0.8333 - val\_loss: 0.5487 - val\_accuracy: 0.7091 - val\_auc: 0.8128  
Epoch 59/100  
2/2 [=====] - 1s 434ms/step - loss: 0.4892 - accuracy:  
0.7778 - auc: 0.8827 - val\_loss: 0.5594 - val\_accuracy: 0.7455 - val\_auc: 0.8135  
Epoch 60/100  
2/2 [=====] - 1s 434ms/step - loss: 0.5419 - accuracy:  
0.6852 - auc: 0.8121 - val\_loss: 0.5539 - val\_accuracy: 0.7273 - val\_auc: 0.8148  
Epoch 61/100  
2/2 [=====] - 1s 499ms/step - loss: 0.4943 - accuracy:  
0.7963 - auc: 0.8683 - val\_loss: 0.5414 - val\_accuracy: 0.7636 - val\_auc: 0.8181  
Epoch 62/100  
2/2 [=====] - 1s 485ms/step - loss: 0.4757 - accuracy:  
0.7963 - auc: 0.8717 - val\_loss: 0.5388 - val\_accuracy: 0.7636 - val\_auc: 0.8188  
Epoch 63/100  
2/2 [=====] - 1s 497ms/step - loss: 0.4740 - accuracy:  
0.7593 - auc: 0.8704 - val\_loss: 0.5376 - val\_accuracy: 0.7636 - val\_auc: 0.8168  
Epoch 64/100  
2/2 [=====] - 1s 437ms/step - loss: 0.4663 - accuracy:  
0.8333 - auc: 0.8923 - val\_loss: 0.5377 - val\_accuracy: 0.7455 - val\_auc: 0.8175  
Epoch 65/100  
2/2 [=====] - 1s 450ms/step - loss: 0.4774 - accuracy:  
0.7593 - auc: 0.8731 - val\_loss: 0.5381 - val\_accuracy: 0.7273 - val\_auc: 0.8181  
Epoch 66/100  
2/2 [=====] - 1s 450ms/step - loss: 0.4645 - accuracy:  
0.8333 - auc: 0.8841 - val\_loss: 0.5410 - val\_accuracy: 0.7636 - val\_auc: 0.8208  
Epoch 67/100  
2/2 [=====] - 1s 467ms/step - loss: 0.4753 - accuracy:  
0.8148 - auc: 0.8813 - val\_loss: 0.5354 - val\_accuracy: 0.7636 - val\_auc: 0.8228  
Epoch 68/100  
2/2 [=====] - 1s 466ms/step - loss: 0.4614 - accuracy:  
0.7593 - auc: 0.8786 - val\_loss: 0.5329 - val\_accuracy: 0.7636 - val\_auc: 0.8241  
Epoch 69/100  
2/2 [=====] - 1s 419ms/step - loss: 0.5030 - accuracy:  
0.7407 - auc: 0.8381 - val\_loss: 0.5334 - val\_accuracy: 0.7636 - val\_auc: 0.8228  
Epoch 70/100  
2/2 [=====] - 1s 487ms/step - loss: 0.4609 - accuracy:  
0.7778 - auc: 0.9053 - val\_loss: 0.5304 - val\_accuracy: 0.7636 - val\_auc: 0.8254  
Epoch 71/100  
2/2 [=====] - 1s 522ms/step - loss: 0.4351 - accuracy:  
0.8148 - auc: 0.9225 - val\_loss: 0.5320 - val\_accuracy: 0.8000 - val\_auc: 0.8228  
Epoch 72/100

2/2 [=====] - 1s 533ms/step - loss: 0.4428 - accuracy:  
0.8333 - auc: 0.9198 - val\_loss: 0.5386 - val\_accuracy: 0.7455 - val\_auc: 0.8254  
Epoch 73/100

2/2 [=====] - 1s 500ms/step - loss: 0.4525 - accuracy:  
0.8519 - auc: 0.9053 - val\_loss: 0.5306 - val\_accuracy: 0.7818 - val\_auc: 0.8254  
Epoch 74/100

2/2 [=====] - 1s 550ms/step - loss: 0.4391 - accuracy:  
0.8148 - auc: 0.9047 - val\_loss: 0.5255 - val\_accuracy: 0.7636 - val\_auc: 0.8267  
Epoch 75/100

2/2 [=====] - 1s 500ms/step - loss: 0.4147 - accuracy:  
0.8704 - auc: 0.9396 - val\_loss: 0.5265 - val\_accuracy: 0.7636 - val\_auc: 0.8261  
Epoch 76/100

2/2 [=====] - 1s 569ms/step - loss: 0.4275 - accuracy:  
0.8519 - auc: 0.9294 - val\_loss: 0.5246 - val\_accuracy: 0.7636 - val\_auc: 0.8307  
Epoch 77/100

2/2 [=====] - 1s 503ms/step - loss: 0.3799 - accuracy:  
0.9259 - auc: 0.9479 - val\_loss: 0.5247 - val\_accuracy: 0.8000 - val\_auc: 0.8307  
Epoch 78/100

2/2 [=====] - 1s 500ms/step - loss: 0.3876 - accuracy:  
0.9074 - auc: 0.9588 - val\_loss: 0.5255 - val\_accuracy: 0.8000 - val\_auc: 0.8320  
Epoch 79/100

2/2 [=====] - 1s 568ms/step - loss: 0.3929 - accuracy:  
0.9259 - auc: 0.9568 - val\_loss: 0.5208 - val\_accuracy: 0.8182 - val\_auc: 0.8333  
Epoch 80/100

2/2 [=====] - 1s 583ms/step - loss: 0.4260 - accuracy:  
0.8704 - auc: 0.9335 - val\_loss: 0.5171 - val\_accuracy: 0.8000 - val\_auc: 0.8366  
Epoch 81/100

2/2 [=====] - 1s 566ms/step - loss: 0.3549 - accuracy:  
0.8889 - auc: 0.9609 - val\_loss: 0.5167 - val\_accuracy: 0.8364 - val\_auc: 0.8360  
Epoch 82/100

2/2 [=====] - 1s 565ms/step - loss: 0.3912 - accuracy:  
0.8704 - auc: 0.9492 - val\_loss: 0.5112 - val\_accuracy: 0.8000 - val\_auc: 0.8353  
Epoch 83/100

2/2 [=====] - 1s 549ms/step - loss: 0.4026 - accuracy:  
0.8519 - auc: 0.9369 - val\_loss: 0.5093 - val\_accuracy: 0.8000 - val\_auc: 0.8327  
Epoch 84/100

2/2 [=====] - 1s 488ms/step - loss: 0.4066 - accuracy:  
0.7963 - auc: 0.9218 - val\_loss: 0.5094 - val\_accuracy: 0.8000 - val\_auc: 0.8287  
Epoch 85/100

2/2 [=====] - 1s 496ms/step - loss: 0.3730 - accuracy:  
0.9074 - auc: 0.9616 - val\_loss: 0.5102 - val\_accuracy: 0.8000 - val\_auc: 0.8287  
Epoch 86/100

2/2 [=====] - 1s 489ms/step - loss: 0.3786 - accuracy:  
0.8889 - auc: 0.9616 - val\_loss: 0.5113 - val\_accuracy: 0.8182 - val\_auc: 0.8280  
Epoch 87/100

2/2 [=====] - 1s 549ms/step - loss: 0.3551 - accuracy:  
0.9259 - auc: 0.9698 - val\_loss: 0.5068 - val\_accuracy: 0.8000 - val\_auc: 0.8307  
Epoch 88/100

2/2 [=====] - 1s 666ms/step - loss: 0.3493 - accuracy:  
0.9259 - auc: 0.9534 - val\_loss: 0.5048 - val\_accuracy: 0.8000 - val\_auc: 0.8300  
Epoch 89/100

2/2 [=====] - 1s 606ms/step - loss: 0.3297 - accuracy:  
0.9259 - auc: 0.9794 - val\_loss: 0.5031 - val\_accuracy: 0.8000 - val\_auc: 0.8313  
Epoch 90/100

2/2 [=====] - 1s 546ms/step - loss: 0.3358 - accuracy:  
0.9444 - auc: 0.9746 - val\_loss: 0.5053 - val\_accuracy: 0.8182 - val\_auc: 0.8360  
Epoch 91/100

2/2 [=====] - 1s 617ms/step - loss: 0.3427 - accuracy:  
0.9444 - auc: 0.9781 - val\_loss: 0.5000 - val\_accuracy: 0.8182 - val\_auc: 0.8300  
Epoch 92/100

2/2 [=====] - 1s 583ms/step - loss: 0.3537 - accuracy:  
0.8704 - auc: 0.9616 - val\_loss: 0.5017 - val\_accuracy: 0.8000 - val\_auc: 0.8307  
Epoch 93/100

2/2 [=====] - 1s 551ms/step - loss: 0.3643 - accuracy:  
0.8519 - auc: 0.9664 - val\_loss: 0.4975 - val\_accuracy: 0.8000 - val\_auc: 0.8327  
Epoch 94/100

2/2 [=====] - 1s 484ms/step - loss: 0.3254 - accuracy:  
0.9259 - auc: 0.9691 - val\_loss: 0.4999 - val\_accuracy: 0.8182 - val\_auc: 0.8360  
Epoch 95/100

2/2 [=====] - 1s 503ms/step - loss: 0.3161 - accuracy:  
0.9259 - auc: 0.9883 - val\_loss: 0.5008 - val\_accuracy: 0.8182 - val\_auc: 0.8446  
Epoch 96/100


2/2 [=====] - 1s 603ms/step - loss: 0.3249 - accuracy:  
0.8704 - auc: 0.9726 - val\_loss: 0.4948 - val\_accuracy: 0.8182 - val\_auc: 0.8433  
Epoch 97/100

2/2 [=====] - 1s 486ms/step - loss: 0.3207 - accuracy:  
0.9259 - auc: 0.9746 - val\_loss: 0.4949 - val\_accuracy: 0.8182 - val\_auc: 0.8419  
Epoch 98/100


2/2 [=====] - 1s 500ms/step - loss: 0.3335 - accuracy:  
0.8704 - auc: 0.9794 - val\_loss: 0.5015 - val\_accuracy: 0.7818 - val\_auc: 0.8433  
Epoch 99/100

2/2 [=====] - 1s 545ms/step - loss: 0.3177 - accuracy:  
0.9444 - auc: 0.9733 - val\_loss: 0.4928 - val\_accuracy: 0.8182 - val\_auc: 0.8439  
Epoch 100/100

2/2 [=====] - 1s 495ms/step - loss: 0.3089 - accuracy:  
0.9259 - auc: 0.9877 - val\_loss: 0.4960 - val\_accuracy: 0.8182 - val\_auc: 0.8439



Final Test Accuracy: 0.6575 | AUC: 0.6821  
 3/3 [=====] - 0s 67ms/step



```
[ ]: # MULTI-CLASS MODEL
x_all, y_all = [], []
all_birds = []

with data as f:
    for idx, key in enumerate(f.keys()):
        spec = f[key][...].astype(np.float32) / 255.0
        spec = np.transpose(spec, (2, 0, 1))
        x_all.append(spec)
        y_all.append(np.full(spec.shape[0], idx, dtype=np.uint8))
        all_birds.append(key)

x_class = np.concatenate(x_all, axis=0)
y_class = np.concatenate(y_all, axis=0)
x_class = x_class[..., np.newaxis]

x_class, y_class = shuffle(x_class, y_class, random_state=42)

x_one, x_test, y_one, y_test = train_test_split(x_class, y_class, test_size=0.
↪4, stratify=y_class, random_state=42)
x_train, x_valid, y_train, y_valid = train_test_split(x_one, y_one, test_size=0.
↪5, stratify=y_one, random_state=42)

x_multi, x_test, y_multi, y_test = train_test_split(x_class, y_class,
                                                    test_size=0.40,
↪stratify=y_class, random_state=42)

x_train, x_valid, y_train, y_valid = train_test_split(x_multi, y_multi,
                                                    test_size=0.50,
↪stratify=y_multi, random_state=42)

num_class = len(all_birds)
print("Input shape: ", x_train.shape[1:], " | classes:", num_class)
```

Input shape: (128, 517, 1) | classes: 12

```
[18]: model_two = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, activation='relu', input_shape=x_train.
↪shape[1:]),
    tf.keras.layers.MaxPooling2D(2),

    tf.keras.layers.Conv2D(64, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(2),

    tf.keras.layers.Conv2D(128, 3, activation='relu'),
    tf.keras.layers.MaxPooling2D(2),
```

```

        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(128, activation='relu'),
        tf.keras.layers.Dense(num_classes, activation='softmax')
    ])

model_two.compile(
    optimizer=tf.keras.optimizers.Adam(1e-4),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

model_two.summary()

callbacks_two = [
    tf.keras.callbacks.EarlyStopping(patience=10, restore_best_weights=True),
    tf.keras.callbacks.ModelCheckpoint('multiclass_birds.keras',
    ↪save_best_only=True)
]

history = model_two.fit(
    x_train, y_train,
    validation_data=(x_valid, y_valid),
    epochs=150,
    batch_size=32,
    callbacks=callbacks_two,
    verbose=1
)

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.legend()

plt.tight_layout()
plt.show()

```



```

loss, acc = model_two.evaluate(x_test, y_test, verbose=0)
print(f"\nFinal Test Accuracy: {acc:.4f}")

probability_two = model_two.predict(x_test)
prediction_two = np.argmax(probability_two , axis=1)
conf_mat = confusion_matrix(y_test, prediction_two)

ConfusionMatrixDisplay(conf_mat, display_labels=all_birds).plot(
    xticks_rotation=90, cmap=plt.cm.Blues)
plt.title('Multiclass Confusion Matrix')
plt.tight_layout()
plt.show()

```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 126, 515, 32)	320
max_pooling2d_18 (MaxPoolin g2D)	(None, 63, 257, 32)	0
conv2d_19 (Conv2D)	(None, 61, 255, 64)	18496
max_pooling2d_19 (MaxPoolin g2D)	(None, 30, 127, 64)	0
conv2d_20 (Conv2D)	(None, 28, 125, 128)	73856
max_pooling2d_20 (MaxPoolin g2D)	(None, 14, 62, 128)	0
flatten_8 (Flatten)	(None, 111104)	0
dense_16 (Dense)	(None, 128)	14221440
dense_17 (Dense)	(None, 12)	1548

```

=====
Total params: 14,315,660
Trainable params: 14,315,660
Non-trainable params: 0

```


```

-----
Epoch 1/150
19/19 [=====] - 33s 2s/step - loss: 2.2840 - accuracy:
0.3047 - val_loss: 2.2106 - val_accuracy: 0.3182
Epoch 2/150
19/19 [=====] - 54s 3s/step - loss: 2.1832 - accuracy:

```

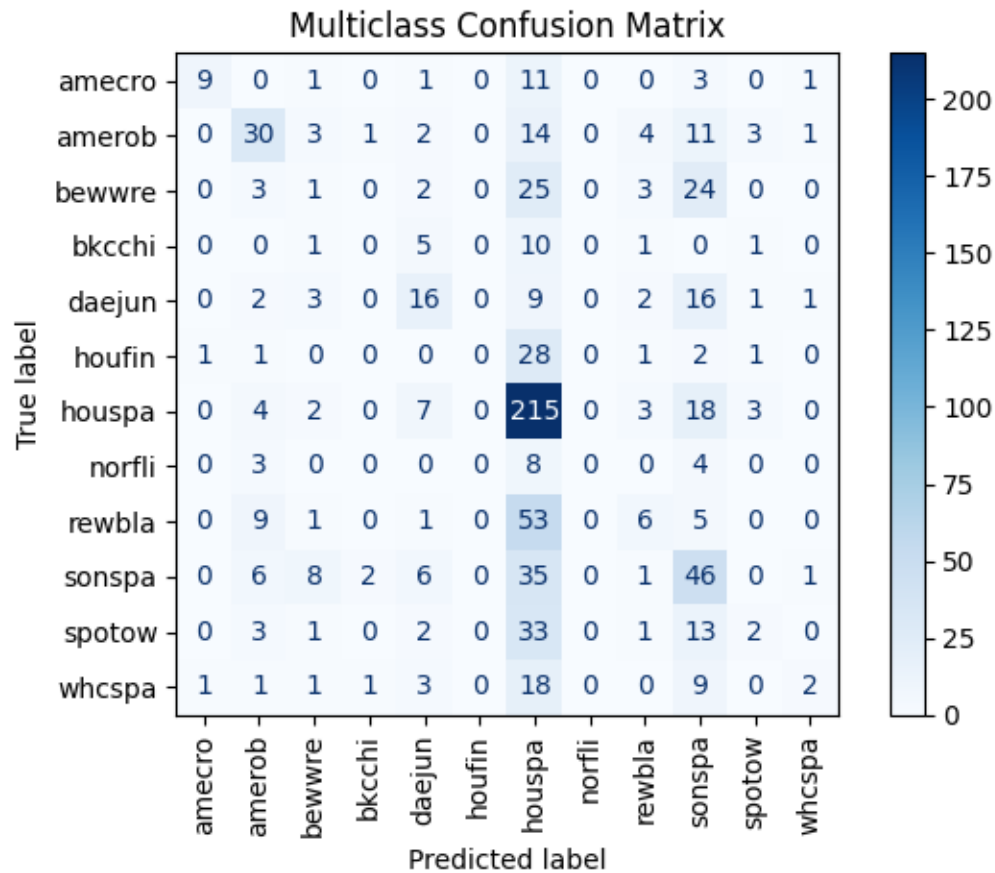
0.3182 - val\_loss: 2.1773 - val\_accuracy: 0.3182  
Epoch 3/150  
19/19 [=====] - 60s 3s/step - loss: 2.1791 - accuracy:  
0.3182 - val\_loss: 2.1681 - val\_accuracy: 0.3182  
Epoch 4/150  
19/19 [=====] - 59s 3s/step - loss: 2.1681 - accuracy:  
0.3182 - val\_loss: 2.1595 - val\_accuracy: 0.3182  
Epoch 5/150  
19/19 [=====] - 58s 3s/step - loss: 2.1556 - accuracy:  
0.3182 - val\_loss: 2.1539 - val\_accuracy: 0.3182  
Epoch 6/150  
19/19 [=====] - 52s 3s/step - loss: 2.1351 - accuracy:  
0.3182 - val\_loss: 2.1412 - val\_accuracy: 0.3182  
Epoch 7/150  
19/19 [=====] - 46s 2s/step - loss: 2.1172 - accuracy:  
0.3182 - val\_loss: 2.1486 - val\_accuracy: 0.3182  
Epoch 8/150  
19/19 [=====] - 33s 2s/step - loss: 2.0828 - accuracy:  
0.3232 - val\_loss: 2.1106 - val\_accuracy: 0.3199  
Epoch 9/150  
19/19 [=====] - 33s 2s/step - loss: 2.0490 - accuracy:  
0.3283 - val\_loss: 2.0888 - val\_accuracy: 0.3249  
Epoch 10/150  
19/19 [=====] - 33s 2s/step - loss: 1.9917 - accuracy:  
0.3384 - val\_loss: 2.0950 - val\_accuracy: 0.3232  
Epoch 11/150  
19/19 [=====] - 41s 2s/step - loss: 1.9549 - accuracy:  
0.3569 - val\_loss: 2.0502 - val\_accuracy: 0.3333  
Epoch 12/150  
19/19 [=====] - 34s 2s/step - loss: 1.8866 - accuracy:  
0.3838 - val\_loss: 2.0182 - val\_accuracy: 0.3737  
Epoch 13/150  
19/19 [=====] - 33s 2s/step - loss: 1.8133 - accuracy:  
0.4158 - val\_loss: 2.0049 - val\_accuracy: 0.3687  
Epoch 14/150  
19/19 [=====] - 33s 2s/step - loss: 1.7176 - accuracy:  
0.4461 - val\_loss: 1.9768 - val\_accuracy: 0.3805  
Epoch 15/150  
19/19 [=====] - 36s 2s/step - loss: 1.6323 - accuracy:  
0.4714 - val\_loss: 1.9693 - val\_accuracy: 0.3906  
Epoch 16/150  
19/19 [=====] - 32s 2s/step - loss: 1.5748 - accuracy:  
0.4916 - val\_loss: 1.9797 - val\_accuracy: 0.3872  
Epoch 17/150  
19/19 [=====] - 32s 2s/step - loss: 1.4796 - accuracy:  
0.5084 - val\_loss: 2.0024 - val\_accuracy: 0.3923  
Epoch 18/150  
19/19 [=====] - 40s 2s/step - loss: 1.4090 - accuracy:

0.5236 - val\_loss: 1.9506 - val\_accuracy: 0.3822  
Epoch 19/150  
19/19 [=====] - 49s 3s/step - loss: 1.3052 - accuracy: 0.5842 - val\_loss: 1.9886 - val\_accuracy: 0.3906  
Epoch 20/150  
19/19 [=====] - 33s 2s/step - loss: 1.2330 - accuracy: 0.6178 - val\_loss: 2.0165 - val\_accuracy: 0.3788  
Epoch 21/150  
19/19 [=====] - 29s 2s/step - loss: 1.1900 - accuracy: 0.6246 - val\_loss: 2.0711 - val\_accuracy: 0.3838  
Epoch 22/150  
19/19 [=====] - 22s 1s/step - loss: 1.1034 - accuracy: 0.6431 - val\_loss: 2.1433 - val\_accuracy: 0.3502  
Epoch 23/150  
19/19 [=====] - 22s 1s/step - loss: 1.0069 - accuracy: 0.6953 - val\_loss: 2.1644 - val\_accuracy: 0.3771  
Epoch 24/150  
19/19 [=====] - 23s 1s/step - loss: 0.9852 - accuracy: 0.6936 - val\_loss: 2.2003 - val\_accuracy: 0.3535  
Epoch 25/150  
19/19 [=====] - 23s 1s/step - loss: 0.9097 - accuracy: 0.7256 - val\_loss: 2.2995 - val\_accuracy: 0.3519  
Epoch 26/150  
19/19 [=====] - 23s 1s/step - loss: 0.8314 - accuracy: 0.7407 - val\_loss: 2.3097 - val\_accuracy: 0.3636  
Epoch 27/150  
19/19 [=====] - 23s 1s/step - loss: 0.7382 - accuracy: 0.7929 - val\_loss: 2.2911 - val\_accuracy: 0.3872  
Epoch 28/150  
19/19 [=====] - 31s 2s/step - loss: 0.6892 - accuracy: 0.8081 - val\_loss: 2.3744 - val\_accuracy: 0.3704



Final Test Accuracy: 0.4124

25/25 [=====] - 9s 359ms/step



```
[ ]: # EXTERNAL TEST DATA
audios = {
    "test1": "/Users/alekh/Desktop/birds/test1.wav",
    "test2": "/Users/alekh/Desktop/birds/test2.wav",
    "test3": "/Users/alekh/Desktop/birds/test3.wav"
}


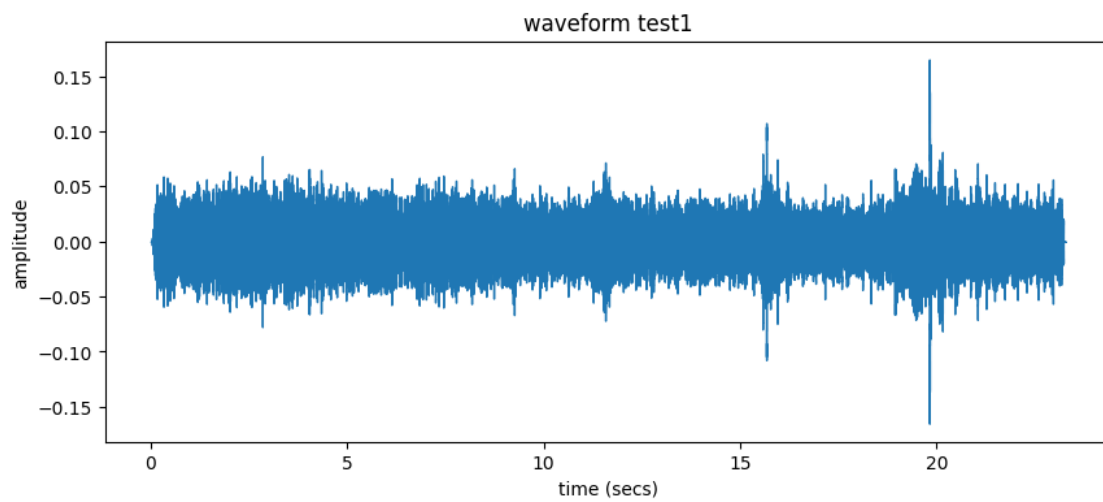
def waveforms(path, name):
    audio, sr = librosa.load(path, sr=22050)


    plt.figure(figsize=(10, 4))
    librosa.display.waveshow(audio, sr=sr)
```

```
plt.title(f"waveform {name}")
plt.xlabel("time (secs)")
plt.ylabel("amplitude")
plt.show()
```

```
return audio, sr
```

```
audio_list = {}
for name, path in audios.items():
    audio, sr = waveforms(path, name)
    audio_list[name] = (audio, sr)
```






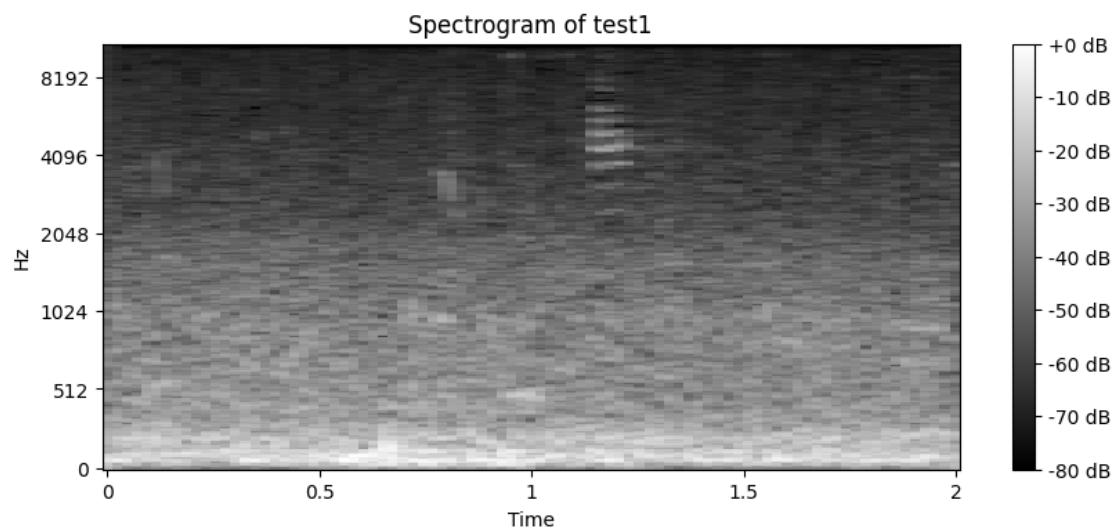
```
[13]: def generate_spectrograms(audio, sr, start_sec, end_sec, name):
    start_sample = int(start_sec * sr)
    end_sample = int(end_sec * sr)
    window = audio[start_sample:end_sample]


    S = librosa.feature.melspectrogram(y=window, sr=sr, n_fft=2048,
    ↪hop_length=512, n_mels=256)
    S_DB = librosa.power_to_db(S, ref=np.max)

    plt.figure(figsize=(10, 4))
    librosa.display.specshow(S_DB, sr=sr, hop_length=512, x_axis='time',
    ↪y_axis='mel', cmap= 'gray')
    plt.colorbar(format='%+2.0f dB')
    plt.title(f'Spectrogram of {name}')
    plt.show()

time_windows = {
    "test1": (15, 17),
    "test2": (0, 2),
    "test3": (2, 4)
}

for name, (audio, sr) in audio_list.items():
    start_sec, end_sec = time_windows[name]
    generate_spectrograms(audio, sr, start_sec, end_sec, name)
```

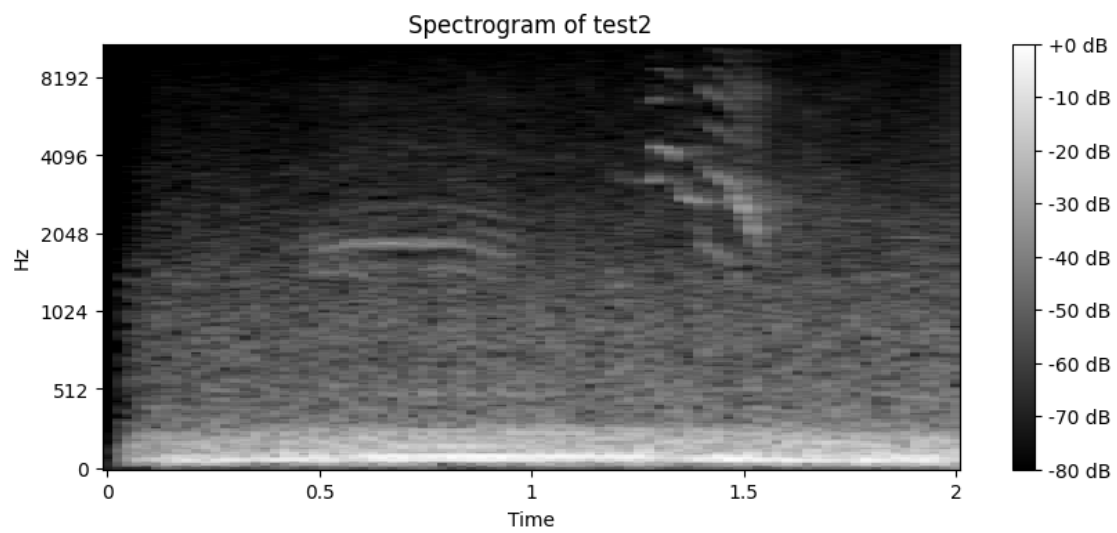






```
[14]: def spectro_array(audio, sr, start_sec, end_sec, n_mels=256, n_fft=2048,
    ↪hop_length=512):
    start_sample = int(start_sec * sr)
    end_sample = int(end_sec * sr)
    window = audio[start_sample:end_sample]
    S = librosa.feature.melspectrogram(y=window, sr=sr, n_fft=n_fft,
    ↪hop_length=hop_length, n_mels=n_mels)
    S_DB = librosa.power_to_db(S, ref=np.max)
    return S_DB

with h5py.File('/Users/alekh/Desktop/birds/spectrograms.h5', 'w') as hf:
    for name, (audio, sr) in audio_list.items():
        start_sec, end_sec = time_windows[name]
        S_DB = spectro_array(audio, sr, start_sec, end_sec)
        hf.create_dataset(name, data=S_DB)
        plt.figure(figsize=(10, 4))
        librosa.display.specshow(S_DB, sr=sr, hop_length=512, x_axis='time',
    ↪y_axis='mel', cmap='grey')
        plt.colorbar(format='%+2.0f dB')
        plt.title(f'Spectrogram of {name}')
        plt.show()
```







```
[16]: import os
import h5py
import numpy as np
import librosa
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.utils import to_categorical
from sklearn.preprocessing import LabelEncoder

spectro_path = '/Users/alekh/Desktop/birds/bird_spectrograms.hdf5'
audios = {
    "test1": "/Users/alekh/Desktop/birds/test1.wav",
    "test2": "/Users/alekh/Desktop/birds/test2.wav",
    "test3": "/Users/alekh/Desktop/birds/test3.wav"
}

time_windows = {
    "test1": (15, 17),
    "test2": (0, 2),
    "test3": (2, 4)
}

spectrogram_list = []
labels = []
with h5py.File(spectro_path, 'r') as f:
    species = list(f.keys())
    for key in species:
        data = f[key][...]
```

```

        if data.ndim == 3:
            for i in range(data.shape[2]):
                spectrogram_list.append(data[:, :, i])
                labels.append(key)
        elif data.ndim == 2:
            spectrogram_list.append(data)
            labels.append(key)
        else:
            raise ValueError(f"ndim={data.ndim}")

spectrogram_list = np.array(spectrogram_list)
labels = np.array(labels)
label_encoder = LabelEncoder().fit(labels)
labels_encoded = label_encoder.transform(labels)
labels_onehot = to_categorical(labels_encoded)
mel_bins, time_frames = spectrogram_list.shape[1], spectrogram_list.shape[2]
spectrogram_list = spectrogram_list.reshape(-1, mel_bins, time_frames, 1).
    ↳astype(np.float32) / 255.0

model_three = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(mel_bins, time_frames,
    ↳1)),
    MaxPooling2D((2,2)),
    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D((2,2)),
    Conv2D(128, (3,3), activation='relu'),
    MaxPooling2D((2,2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(len(species), activation='softmax')
])
model_three.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])
model_three.fit(spectrogram_list, labels_onehot, epochs=5, batch_size=32)

test_data = []
for name, path in audios.items():
    audio, sr = librosa.load(path, sr=22050)
    start, end = time_windows[name]
    clip = audio[int(start*sr):int(end*sr)]
    S = librosa.feature.melspectrogram(y=clip, sr=sr, n_fft=2048,
    ↳hop_length=512, n_mels=mel_bins)
    S_db = librosa.power_to_db(S, ref=np.max)
    if S_db.shape[1] < time_frames:
        S_db = np.pad(S_db, ((0,0),(0, time_frames - S_db.shape[1])),
    ↳'constant')

```

```

else:
    S_db = S_db[:, :time_frames]
    test_data.append(S_db)

test_data = np.array(test_data).reshape(-1, mel_bins, time_frames, 1).astype(np.
    ↪float32) / 255.0
test_predictions = model_three.predict(test_data, verbose=0)
test_one = np.argmax(test_predictions, axis=1)
test_class = label_encoder.inverse_transform(test_one)
for i, pred in enumerate(test_class, 1):
    print(f"Test Spectrogram {i} predicted as {pred}")

```

```

Epoch 1/5
62/62 [=====] - 73s 1s/step - loss: 2.2660 - accuracy:
0.2968
Epoch 2/5
62/62 [=====] - 63s 1s/step - loss: 2.2259 - accuracy:
0.3019
Epoch 3/5
62/62 [=====] - 65s 1s/step - loss: 2.2022 - accuracy:
0.3125
Epoch 4/5
62/62 [=====] - 65s 1s/step - loss: 2.1840 - accuracy:
0.3150
Epoch 5/5
62/62 [=====] - 66s 1s/step - loss: 2.1032 - accuracy:
0.3241
Test Spectrogram 1 predicted as sonspa
Test Spectrogram 2 predicted as sonspa
Test Spectrogram 3 predicted as sonspa

```